

Tutorials 教程Intro to LangGraph LangGraph
简介

Use cases 用例
 Chatbots 聊天机器人 >
 Multi-Agent Systems 多代理系统
 RAG >
 Web Research (STORM) 网络研究 (风暴)
 Planning Agents 规划代理 >
Plan-and-Execute 计划与执行
 Reasoning w/o Observation 不观察的推理
 LLMCompiler 法学硕士编译器
 Reflection & Critique 反思与批评
 Evaluation & Analysis 评估 >
 与分析
 Web Navigation 网页导航
 Competitive Programming 竞争性编程

Plan-and-Execute 计划与执行 ¶

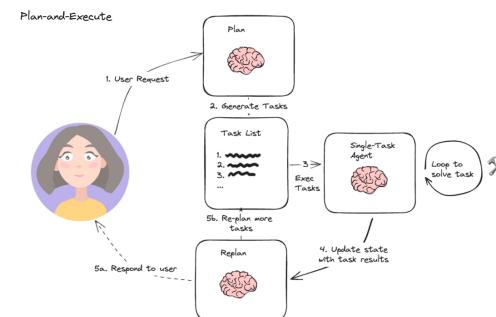
This notebook shows how to create a "plan-and-execute" style agent. This is heavily inspired by the [Plan-and-Solve](#) paper as well as the [Baby-AGI](#) project. 本笔记本展示了如何创建"计划并执行"风格的代理。这很大程度上受到了 Plan-and-Solve 论文以及 Baby-AGI 项目的启发。

The core idea is to first come up with a multi-step plan, and then go through that plan one item at a time. After accomplishing a particular task, you can then revisit the plan and modify as appropriate.

核心思想是首先制定一个多步骤计划，然后一次一项地执行该计划。完成特定任务后，您可以重新审视计划并进行适当修改。

The general computational graph looks like the following:

一般的计算图如下所示：



This compares to a typical [ReAct](#) style agent where you think one step at a time. The advantages of this "plan-and-execute" style agent are:

这与典型的 ReAct 风格代理相比，您一次思考一步。这种"计划和执行"风格代理的优点是：

1. Explicit long term planning (which even really strong LLMs can struggle with)
明确的长期规划（即使是非常强大的LLMs也可能遇到困难）
2. Ability to use smaller/weaker models for the execution step, only using larger/better models for the planning step
能够在执行步骤中使用较小/较弱的模型，仅在规划步骤中使用较大/更好的模型

The following walkthrough demonstrates how to do so in LangGraph. The resulting agent will leave a trace like the following example: ([link](#)).

以下演练演示了如何在 LangGraph 中执行此操作。生成的代理将留下如下例所示的痕迹：（链接）。

Setup 设置

First, we need to install the packages required.

首先，我们需要安装所需的软件包。

```
In [1]: 在 [1] 中: %capture --no-stderr
!pip install --quiet -U langchain==community langchain==openai tavyly==pyt
```

Next, we need to set API keys for OpenAI (the LLM we will use) and Tavyly (the search tool we will use)

接下来，我们需要为 OpenAI（我们将使用的 LLM）和 Tavyly（我们将使用的搜索工具）设置 API 密钥

```
In [1]: 在 [1] 中:
import os
import getpass

def _set_env(var: str):
    if not os.environ.get(var):
        os.environ[var] = getpass.getpass(f"({var}): ")

_set_env("OPENAI_API_KEY")
_set_env("TAVLY_API_KEY")
```

Optionally, we can set API key for LangSmith tracing, which will give us best-in-class observability.

或者，我们可以为 LangSmith 跟踪设置 API 密钥，这将为我们提供一流的可观测性。

```
In [2]: 在 [2] 中:
os.environ["LANGCHAIN_TRACING_V2"] = "true"
_set_env("LANGCHAIN_API_KEY")
os.environ["LANGCHAIN_PROJECT"] = "Plan-and-execute"
```

Define Tools 定义工具

We will first define the tools we want to use. For this simple example, we will use a built-in search tool via Tavyly. However, it is really easy to create your own

Table of contents

目录

Setup 设置

Define Tools 定义工具

Define our Execution Agent 定义我们的执行代理

Define the State 定义状态

Planning Step 规划步骤

Re-Plan Step 重新计划步骤

Create the Graph 创建图表

Conclusion 结论

tools - see documentation [here](#) on how to do that.

我们将首先定义我们想要使用的工具。对于这个简单的示例，我们将通过 Tavily 使用内置搜索工具。然而，创建自己的工具确实很容易 - 请参阅此处的文档了解如何做到这一点。

```
In [3]: 在 [3] 中:   from langchain_community.tools.tavily_search import TavilySearchResults  
                           tools = [TavilySearchResults(max_results=3)]
```

Define our Execution Agent 定义我们的执行代理

Now we will create the execution agent we want to use to execute tasks. Note that for this example, we will be using the same execution agent for each task but this doesn't HAVE to be the case.

现在我们将创建要用来执行任务的执行代理。请注意，对于此示例，我们将为每个任务使用相同的执行代理，但情况并非必须如此。

```
In [5]: In [5]: from langchain import hub
from langchain_openai import ChatOpenAI
from langgraph.prebuilt import create_react_agent

# Get the prompt to use - you can modify this!
prompt = hub.pull("wfh/react-agent-executor")
prompt.pretty_print()

# Choose the LLM that will drive the agent
llm = ChatOpenAI(model="gpt-4-turbo-preview")
agent_executor = create_react_agent(llm, tools, messages_modifier=prefix)

=====
***** System Message *****
You are a helpful assistant.

*****
***** Messages Placeholder *****
{{messages}}
```

Define the State 定义状态

Let's now start by defining the state the track for this agent.
现在让我们开始定义该代理的跟踪状态。

First, we will need to track the current plan. Let's represent that as a list of strings.

首先，我们需要跟踪当前的计划。让我们将其表示为字符串列表。

Next, we should track previously executed steps. Let's represent that as a list of tuples (these tuples will contain the step and then the result)

接下来，我们应该跟踪之前执行的步骤。让我们将其表示为元组列表（这些元组将包含步骤，然后包含结果）

Finally, we need to have some state to represent the

最后，我们需要一些状态来表示最终响应以及原始输入。

```
In [8]: 在[8]中: from typing import List, Tuple, Annotated, Type  
import operator
```

```
import operator
```

```

class PlanExecute(TypedDict):
    input: str
    plan: List[str]
    past_steps: Annotated[List[Tuple], operator.add]
    response: str

```

Planning Step 规划步骤

Let's now think about creating the planning step. This will use function calling to create a plan.

现在让我们考虑创建计划步骤。这将使用函数调用来创建计划。

```

In [10]: 在 [10] 中: from langchain_core.pydantic_v1 import BaseModel, Field

class Plan(BaseModel):
    """Plan to follow in future"""

    steps: List[str] = Field(
        description="different steps to follow, should be in sorted order"
    )

In [36]: 在 [36] 中: from langchain_core.prompts import ChatPromptTemplate

planner_prompt = ChatPromptTemplate.from_messages(
    [
        (
            "system",
            """For the given objective, come up with a simple step by step plan. This plan should involve individual tasks, that if executed correctly, the result of the final step should be the final answer. Make sure the steps are in sorted order.
            """,
            "placeholder", {"messages": ""}
        )
    ]
)
planner = planner_prompt | ChatOpenAI(
    model="gpt-4o", temperature=0
).with_structured_output(Plan)

In [37]: 在 [37] 中: planner.invoke(
    {
        "messages": [
            ("user", "what is the hometown of the current Australian Open winner?")
        ]
    }
)

Out[37]: 输出[37]: Plan(steps=['Identify the current winner of the Australian Open.', 'Determine the hometown of the identified winner.'])

```

Re-Plan Step 重新计划步骤

Now, let's create a step that re-does the plan based on the result of the previous step.

现在，让我们创建一个步骤，根据上一步的结果重新执行计划。

```

In [19]: 在 [19] 中: from typing import Union

class Response(BaseModel):
    """Response to user"""

    response: str

class Act(BaseModel):
    """Action to perform"""

    action: Union[Response, Plan] = Field(
        description="Action to perform. If you want to respond to user, use Response. If you need to further use tools to get the answer, use Plan."
    )

replanner_prompt = ChatPromptTemplate.from_template(
    """For the given objective, come up with a simple step by step plan. This plan should involve individual tasks, that if executed correctly, the result of the final step should be the final answer. Make sure the steps are in sorted order.
    """
)
Your objective was this:
{input}
Your original plan was this:
{plan}
You have currently done the following steps:
{past_steps}
Update your plan accordingly. If no more steps are needed and you can stop, just return.
)

replanner = replanner_prompt | ChatOpenAI(
    model="gpt-4o", temperature=0
).with_structured_output(Act)

```

Create the Graph 创建图表

We can now create the graph!

我们现在可以创建图表了！

```

In [54]: 在 [54] 中: from typing import Literal

async def execute_step(state: PlanExecute):
    plan = state["plan"]
    plan_str = "\n".join(f"{i+1}. {step}" for i, step in enumerate(plan))
    task = plan[0]
    task_formatted = f"For the following plan:\n{plan_str}\n\nYou are tasked with executing step {1}, ({task})."
    agent_response = await agent_executor.ainvoke(
        {"messages": [{"user": task_formatted}]}
    )
    return {
        "past_steps": (task, agent_response["messages"][-1].content),
        "plan": plan[1:]
    }

async def plan_step(state: PlanExecute):
    plan = await replanner.ainvoke({"messages": [{"user": state["input"]}]}
    return {"plan": plan.steps}

async def replan_step(state: PlanExecute):
    output = await replanner.ainvoke(state)
    if isinstance(output.action, Response):
        return {"response": output.action.response}
    else:
        return {"plan": output.action.state}

```

```

def should_end(state: PlanExecute) -> Literal["agent", "__end__"]:
    if "response" in state and state["response"]:
        return "__end__"
    else:
        return "agent"

In [55]: In [55]: from langgraph.graph import StateGraph
workflow = StateGraph(PlanExecute)

# Add the plan node
workflow.add_node("planner", plan_step)

# Add the execution step
workflow.add_node("agent", execute_step)

# Add a replan node
workflow.add_node("replan", replan_step)

workflow.set_entry_point("planner")

# From plan we go to agent
workflow.add_edge("planner", "agent")

# From agent, we replan
workflow.add_edge("agent", "replan")

workflow.add_conditional_edges(
    "replan",
    # Next, we pass in the function that will determine which node is
    # called next based on the state
    should_end,
)

```

Finally, we compile it!
This compiles it into a LangChain Runnable,
meaning you can use it as you would any other runnable
app = workflow.compile()

```

In [56]: In [56]: from IPython.display import Image, display
display(Image(app.get_graph(xray=True).draw_mermaid_png()))

```



```

In [57]: In [57]: config = {"recursion_limit": 50}
inputs = {"input": "what is the hometown of the 2024 Australia open winner?"}
async for event in app.astream(inputs, config=config):
    for k, v in event.items():
        if k != "__end__":
            print(v)

```

{'plan': ['Identify the winner of the 2024 Australian Open.', 'Determine the hometown of the identified winner.'],
'past_steps': ['Identify the winner of the 2024 Australian Open.', 'The winner of the 2024 Australian Open is Jannik Sinner. He claimed his first Grand Slam title in an epic comeback win over Daniil Medvedev.'],
'plan': ['Determine the hometown of Jannik Sinner.'],
'past_steps': ['Determine the hometown of Jannik Sinner.', 'Jannik Sinner's hometown is not directly mentioned in the provided excerpts. To ensure accurate information, it's advisable to check a reliable source like his official ATP profile or a detailed biography which often includes personal background details such as hometown.'],
'plan': ['Check Jannik Sinner's official ATP profile or a detailed biography to find his hometown.', 'Return the hometown of Jannik Sinner.'],
'past_steps': ['Check Jannik Sinner's official ATP profile or a detailed biography to find his hometown.', 'Jannik Sinner's official ATP profile can be found at this URL: [ATP Tour - Jannik Sinner](https://www.atptour.com/en/players/jannik-sinner/seag/overview). This profile will contain detailed information including his biography, rankings, playing activity, and potentially his hometown.'],
'plan': ['Visit Jannik Sinner's official ATP profile or a detailed biography to find his hometown.', 'Return the hometown of Jannik Sinner.'],
'past_steps': ['Visit Jannik Sinner's official ATP profile or a detailed biography to find his hometown.', 'Jannik Sinner's official ATP profile and other reliable sources do not explicitly mention his hometown in the search results provided. For detailed information, visiting his ATP profile directly or consulting a comprehensive biography would be recommended to find this specific information.'],
'plan': ['Visit Jannik Sinner's official ATP profile or a detailed biography to find his hometown.', 'Return the hometown of Jannik Sinner.'],
'past_steps': ['Visit Jannik Sinner's official ATP profile or a detailed biography to find his hometown.', 'Jannik Sinner's official ATP profile can be accessed at [here](https://www.atptour.com/en/players/jannik-sinner/seag/overview), although it does not directly provide his hometown in the snippet. For detailed information, such as his hometown, it might be necessary to visit the profile directly or consult other detailed biographies like the one available on [Wikipedia](http://en.wikipedia.org/wiki/Jannik_Sinner), which often include personal details such as hometowns.'],
'plan': ['Visit Jannik Sinner's official ATP profile or his Wikipedia page to find his hometown.', 'Return the hometown of Jannik Sinner.'],
'past_steps': ['Visit Jannik Sinner's official ATP profile or his Wikipedia page to find his hometown.', 'Jannik Sinner's official ATP profile and Wikipedia page did not directly mention his hometown in the provided excerpts. However, further information can typically be found by visiting the full pages directly through the provided links: [ATP Tour - Jannik Sinner's ATP Tour Profile](https://www.atptour.com/en/players/jannik-sinner/seag/overview) and [Jannik Sinner's Wikipedia Page](http://en.wikipedia.org/wiki/Jannik_Sinner). These pages likely contain detailed information, including his hometown. I recommend checking these sources.'],
'response': 'The necessary steps to find the hometown of the 2024 Australian Open winner, Jannik Sinner, have already been completed. His hometown is Immenchen, Italy.'}

Conclusion 结论

Congrats on making a plan-and-execute agent! One known limitation of the above design is that each task is still executed in sequence, meaning embarrassingly parallel operations all add to the total execution time. You could improve on this by having each task represented as a DAG (similar to LLMCompiler), rather than a regular list.

恭喜您成为计划和执行代理！上述设计的一个已知限制是每个任务仍然按顺序执行，这意味着令人尴尬的并行操作都会增加总执行时间。您可以通过将每个任务表示为 DAG（类似于 LLMCompiler）而不是常规列表来改进这一点。

```
In [1]: In [1]:
```

Comments 评论



2 reactions



2

1 comment — powered by giscus

Oldest Newest

 levbszabo 4 days ago

Curious if anyone has ever built anything leveraging this in combination with 'human in the loop' components - meaning AI plans and executes the overall project and escalates to different human agents when needed.

↑ 1 

0 replies

Write Preview Aa

Sign in to comment

 Sign in with GitHub

Previous 以前的
← Web Research (STORM) 网络研究 (风暴)

Next 下一个
Reasoning w/o Observation
不观察的推理 →

Made with Material for McGraw-Hill Education

